


The sparse representation and the edit actually applied

TopK controls support size directly; error-preserving edits change selected decoder contributions without replacing the full MLP activation



ERROR-PRESERVING CAUSAL EDITS

INHIBIT A FROZEN FEATURE SET S

$$x'_\ell = x_\ell - s_\ell W_{d,S} z_{\ell,S}$$

subtract only the selected decoder contribution

SWAP THE SAME SET FROM A DONOR

$$x'_\ell = x_\ell + s_\ell W_{d,S} (z_{\ell,S}^{src} - z_{\ell,S}^{tgt})$$

retain the target reconstruction error and all unedited directions